



中山大學
SUN YAT-SEN UNIVERSITY

Lecture 3

Solving Problems by Searching

Algorithm


张子臻，中山大学计算机学院

zhangzizhen@gmail.com

Outline

- Problem solving agents
- Search examples
- Uninformed search
- Informed search

Agents and Environments



- Agents include humans, robots, softbots, thermostats, etc.
- The agent function maps from percept histories to actions:

$$f : \mathcal{P}^* \rightarrow \mathcal{A}$$

Agents

- **Simple-reflex agents** directly maps states to actions. They may not operate well in environments where the mapping is too large to store or takes too much to learn.
- **Goal-based agents** can succeed by considering future actions and desirability of their outcomes.
- **Problem solving agent** is a goal-based agent that decides what to do by finding sequences of actions that lead to desirable states.
 - Looking for such a sequence is called **search**.
 - A search algorithm takes a problem as input and returns a solution in the form of **action sequence**.
 - Once a solution is found the actions it recommends can be carried out – **execution** phase.

Components of the Search

- A search problem can be defined formally by four components.
 - Initial state
 - Goal test
 - Successor function
 - Path cost function that assigns a numeric cost to each path. The cost of a path can be described as the sum of the costs of the individual actions along the path – step cost
- Problem solution: sequences of actions to be taken successively.

Search Example: 8-Puzzle Problem

7	2	4
5		6
8	3	1

Start State


1	2	3
4	5	6
7	8	

Goal State

- States: locations of tiles
- Initial state: start state (given)
- Goal test: goal state (given)
- Successor function
 - Action: move blank left, right, up, down
- Path cost: 1 per move


State vs. Node

- A state is a (representation of) a physical configuration.
- A node is a data structure constituting part of a search tree includes **state**, **parent node**, **action**, **path cost $g(x)$** , **depth**.




Search Example: River Crossing

- State space **S**: all valid configurations
- Initial states (nodes) $I = \{(CSDF,)\} \subseteq S$
 - Where's the boat?
- Goal states $G = \{(\,,CSDF)\} \subseteq S$
- Successor function $\text{succs}(s) \subseteq S$: states reachable in one step (one arc) from **S**
 - $\text{succs}((CSDF,)) = \{(CD, SF)\}$
 - $\text{succs}((CDF,S)) = \{(CD,FS), (D,CFS), (C, DFS)\}$
- $\text{cost}(s,s') = 1$ for all arcs. (weighted for other problems)
- The search problem: find a solution path from a state in **I** to a state in **G**.
 - Optionally minimize the cost of the solution.



Search Graph in State Space



- In general, there will be many generated, but un-expanded states at any given time.
- One has to choose which one to expand next.
- Search strategy: a function that selects the next node to be expanded.

Search Strategies

- A search strategy is defined by picking the order of node expansion.
- Strategies are evaluated along the following dimensions:
 - Completeness: does it always find a solution if one exists?
 - Time complexity: number of nodes generated
 - Space complexity: maximum number of nodes in memory
 - Optimality: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - ***b***: maximum branching factor of the search tree
 - ***d***: depth of the least-cost solution
 - ***m***: maximum depth of the state space (may be ∞)

Uninformed Search

- Uninformed means we only know:
 - The goal test
 - The **succs()** function
- But not which non-goal states are better: that would be informed search.
- For now, we also assume **succs()** graph is a **tree**.
 - Won't encounter repeated states.
 - We will discuss it later.
- Search strategies: BFS, UCS, DFS, IDS.


Breadth-First Search (BFS)

- Use a queue (First-in First-out)

```

en_queue(Initial states)
While (queue not empty)
    s = de_queue()
    if (s==goal) success!
    T = succs(s)
    for t in T: t.prev=s
    en_queue(T)
endWhile
  
```

We need back
pointers to recover
the solution path.



Performance of BFS

- Assume:
 - the graph may be infinite.
 - Goal(s) exists and is only finite steps away.
- Will BFS find at least one goal?
- Will BFS find the least cost goal?
- Time complexity?
 - Number of states generated
 - Goal: d edges away
 - Branching factor: b
- Space complexity?
 - Number of states stored

Performance of BFS

- Completeness: yes, BFS will find a goal.
- Optimality: yes, if edges cost 1 (more generally positive non-decreasing in depth), no otherwise.
- Time complexity (worst case): goal is the last node at radius **d** .
 - Have to generate all nodes at radius **d** .
 - **$b + b^2 + \dots + b^d \sim O(b^d)$**
- Space complexity: **$O(b^d)$**

Uniform-Cost Search


- Find the least-cost goal
- Each node has a path cost from start (= sum of edge costs along the path). Expand the least cost node first.
- Use a priority queue instead of a normal queue
 - Always take out the least cost item
 - Remember *heap*? time $O(\log(\text{number of items in heap}))$
- Complete and optimal (if edge costs $\geq \epsilon > 0$)
- Time and space: can be much worse than BFS
 - Let C^* be the cost of the least-cost goal
 - $O(b^{C^*/\epsilon})$, possibly $C^*/\epsilon \gg d$

Depth-First Search

- Use a stack (First-in Last-out)

```
push(Initial states)
While (stack not empty)
    s = pop()
    if (s==goal) success!
    T = succs(s)
    push(T)
endWhile
```

↖
This is non-recursive
implementation of DFS, recursive
implementation is more common




Performance of DFS

- m = maximum depth of graph from start
- Space complexity: $O(mb)$
- Infinite tree: may not find goal (incomplete)
- May not be optimal
- Finite tree: may visit almost all nodes, time complexity $O(b^m)$

Iterative Deepening Search

1. DFS, but stop if path length > 1 .
2. If goal not found, repeat DFS, stop if path length > 2 .
3. And so on...




Performance of IDS

- BFS + DFS
 - Complete, optimal like BFS
 - Small space complexity like DFS
- A huge waste?
 - Each deepening repeats DFS from the beginning
 - No! $db + (d-1)b^2 + (d-2)b^3 + \dots + b^d \sim O(b^d)$
 - Time complexity like BFS

Bidirectional Search

- Breadth-first search from both start and goal
- Stop when fringes meet
- The fringes(边缘) are $O(b^{d/2})$
- Generates $O(b^{d/2})$ instead of $O(b^d)$ nodes




Performance of Search Algorithms

b: branching factor (assume finite) d: goal depth m: graph depth

	Complete	optimal	time	space
Breadth-first search	Y	Y, if ¹	$O(b^d)$	$O(b^d)$
Uniform-cost search	Y	Y	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Depth-first search	N	N	$O(b^m)$	$O(bm)$
Iterative deepening	Y	Y, if ¹	$O(b^d)$	$O(bd)$
Bidirectional search	Y	Y, if ¹	$O(b^{d/2})$	$O(b^{d/2})$

1. edge cost constant, or positive non-decreasing in depth


An Example



- All edges are directed, pointing downwards


Node Expansion

- Depth-First Search:
 - S A D E G
 - Solution found: S A G
- Breadth-First Search:
 - S A B C D E G
 - Solution found: S A G
- Uniform-Cost Search:
 - S A D B C E G
 - Solution found: S B G (This is the only uninformed search that worries about costs.)
- Iterative-Deepening Search:
 - S A B C S A D E G
 - Solution found: S A G



General Graph Search

- The problem: repeated states



- We have to remember already-expanded states (**CLOSED**).
- When we take out a state from the fringe (**OPEN**), check whether it is in **CLOSED** (already expanded).
 - If yes, throw it away.
 - If no, expand it (add successors to **OPEN**), and move it to **CLOSED**.

General Graph Search


- BFS:
 - Still $O(b^d)$ space complexity
- DFS:
 - Memorizing DFS (MEMDFS): memorize every expanded states
 - Path Check DFS (PCDFS): remember only expanded states on current path (from start to the current node)

Informed Search

- Uninformed search
 - Knows the actual path cost $g(s)$ from the start to a node s .




- Informed search
 - Also has a heuristic $h(s)$ of the cost from s to goal.
 - Can be much faster than uninformed search.




Recall: Uniform-Cost Search

- Uniform-cost search: uninformed search when edge costs are not the same.
- Complete (will find a goal) and Optimal (will find the least-cost goal).
- Always expand the node with the least $g(s)$
- Use a priority queue:
 - Push in states with their first-half-cost $g(s)$
 - Pop out the state with the least $g(s)$ first
- Now we have an estimate of the second-half-cost $h(s)$, how to use it?



Best-First Greedy Search


- Use $h(s)$ instead of $g(s)$
- Always expand the node with the least $h(s)$
- Not optimal



It will follow the path $A \rightarrow C \rightarrow G$


A Search

- Use $f(s)=g(s)+h(s)$
- Always expand the node with the least $g(s)+h(s)$
- A search is not always optimal



A* Search

- Same as A search, but the heuristic function $h()$ has to satisfy $h(s) \leq h^*(s)$, where $h^*(s)$ is the true cost from node s to the goal.
- Such heuristic function $h()$ is called **admissible**.
- An admissible heuristic never over-estimates.
- A search with admissible $h()$ is called **A* search**.



Admissible Heuristic Functions h

- 8-puzzle example

Example State

1		5
2	6	3
7	4	8

Goal State

1	2	3
4	5	6
7	8	


- Which of the following are admissible heuristics?
 - $h(n)$ =number of tiles in wrong position
 - $h(n)=0$
 - $h(n)=1$
 - $h(n)$ =sum of Manhattan distance between each tile and its goal location

Admissible Heuristics


- A heuristic function h_2 **dominates** h_1 if for all s
 $h_1(s) \leq h_2(s) \leq h^*(s)$
- $d = 14$,
 - $A^*(h_1) = 539$ nodes
 - $A^*(h_2) = 113$ nodes
- $d = 24$,
 - $A^*(h_1) = 39,135$ nodes
 - $A^*(h_2) = 1,641$ nodes
- We prefer heuristic functions as close to h^* as possible, but not over h^* .
- Good heuristic function might need complex computation.
- Time may be better spent, if we use a faster, simpler heuristic function and expand more nodes.

Some Tricks

- A^* should terminate only when a goal is popped from the priority queue



- A^* can revisit an expanded state, and discover a shorter path



The A* Algorithm

1. Put the start node **S** on the priority queue, called **OPEN**
2. If **OPEN** is empty, exit with failure
3. Remove from **OPEN** and place on **CLOSED** a node **n** for which **f(n)** is minimum
4. If **n** is a goal node, exit (trace back pointers from **n** to **S**)
5. Expand **n**, generating all its successors and attach to them pointers back to **n**.
For each successor **n'** of **n** not on **CLOSED**
 1. If **n'** is not already on **OPEN**, estimate $h(n')$, $g(n')=g(n)+c(n,n')$, $f(n')=g(n')+h(n')$, and place it on **OPEN**.
 2. If **n'** is already on **OPEN**, then check if $g(n')$ is lower for the new version of **n'**. If so, then:
 - Redirect pointers backward from **n'** along path yielding lower $g(n')$.
 - Put **n'** on **OPEN**.
 - If $g(n')$ is not lower for the new version, do nothing.
6. Goto 2.

Use priority queue

Use hashing

Choosing a Hash Function

- A hash function should be easy and quick to compute.
- A hash function should achieve an **even distribution** of the keys that actually occur across the range of indices.
- The usual way to make a hash function is to take the key, chop it up, mix the pieces together in various ways, and thereby obtain an index that will be **uniformly distributed** over the range of indices.

Hashing Functions for Permutations


- Sometimes, when keys are not integers, try to convert them to integers.
- For example, how to convert the permutation (2,4,1,3) to a number?
 - Cantor expansion / reverse Cantor expansion
 - <http://baike.baidu.com/view/437641.htm>

```
int PermutationToNumber(int permutation[], int n) {  
    int result = 0;  
    for (int i = 0; i < n; i++) {  
        int count = 0;  
        for (int j = i + 1; j < n; j++) {  
            if (permutation[j] < permutation[i]) count++;  
        }  
        // factorials[j] records j!  
        result += count * factorials[n - i - 1];  
    }  
    return result;  
}
```

1	2	3
4	5	6
7	8	


Consistent Heuristics

- Consistency is analogous to the triangle inequality from Euclidian geometry.
- A heuristic is **consistent** if
$$h(n) \leq c(n, a, n') + h(n')$$
- If h is consistent, then for every child n' of n , we have:
$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n, a, n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$
- That is, $f(n)$ is non-decreasing along any path.



Behavior of A^* with Consistent Heuristic

- If h is consistent, then A^* expands nodes in order of increasing f value. In such a case, A^* can be implemented more efficiently — no node needs to be processed more than once.
- Gradually adds “ f -contours” of nodes (breadth-first adds layers)
- Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$




Properties of A*

- Complete? Yes
- Time? $O(\text{entire state space})$ in worst case, $O(d)$ in best case
- Space? Keeps all nodes in memory
- Optimal? Yes

A* for Traveling Salesman Problem

- For a node s
$$f(s)=g(s)+h(s)$$
- What is $g(s)$, $h(s)$?
- A* v.s. branch-and-bound



Iterative-Deepening A*

```
function IDA*(problem) returns a solution
  inputs: problem, a problem
   $f_0 \leftarrow h(\text{initial state})$ 
  for  $i \leftarrow 0$  to  $\infty$  do
     $\text{result} \leftarrow \text{COST-LIMITED-SEARCH}(\text{problem}, f_i)$ 
    if  $\text{result}$  is a solution then return  $\text{result}$ 
    else  $f_{i+1} \leftarrow \text{result}$ 
  end


function COST-LIMITED-SEARCH(problem,  $f_{\max}$ ) returns solution or number
  depth-first search, backtracking at every node  $n$  such that  $f(n) > f_{\max}$ 
  if the search finds a solution then
    return the solution
  else
    return  $\min\{f(n) \mid \text{the search backtracked at } n\}$ 
```

Variants of A*

- <https://theory.stanford.edu/~amitp/GameProgramming/Variations.html>

- Beam search
- Iterative deepening
- Weighted A*
- Bandwidth search
- Bidirectional search
- Dynamic A* and Lifelong Planning A*
- ...

- Multi-Agent Path Finding
 - Conflict-Based Search (CBS)



Thank you!

